A Systematic Review of AI-Powered Payment Fraud Detection and Optimization

Krishna K G[1*], Deepak Yadav[1], Deepa S G[1], Shivaprasad Ashok Chikop[1], and Deepa Mathew K[1]

Department of AI & ML, Dayananda Sagar Academy of Technology and Management, Bangalore, India
krishnakg1205@gmail.com* (Corresponding Author)
deepakrises2002@gmail.com
deepasg2110@gmail.com
shivaprasad-aiml@dsatm.edu.in
deepa-aiml@dsatm.edu.in

Abstract. Annual financial losses exceeding $485.6B from global digital payments experienced massive and rapid growth that has magnified the scope and risk of payment fraud. Real-time automated fraud detection and transaction optimization have become possible through Artificial Intelligence and Machine Learning. Using a systematic review of more than 124 peer-reviewed papers published over the last decade, including supervised, unsupervised, anomaly-based, graph-based, deep learning, and hybrid systems used in payment systems, this study examines the most pressing global bank issues, payment providers, and central processing units. Optimum frameworks, theoretical structures, and assessment tactics are defined. With the advent of blockchain and quantum computing, learning. In addition, the trained model analyzed the current and possible problems predicted by academia, the industry , and the implementation of optimized fraud detection strategies. Probably, evaluating and detecting the potential under the paper issue, reducing the automatic calculating fictional total cost responsive, and optimizing the live-routing response may have very important results.

Keywords: AI · Payment Fraud Detection · Machine Learning · Deep Learning · Real-time Processing · Concept Drift · Explainable AI · Optimization · Edge Computing · Federated Learning · Blockchain · Quantum Computing

1 Introduction

The surge of digital financial services and instant payment solutions has led to an increasingly rampant incidence of payment fraud that threatens global economies and consumer welfare. According to Nasdaq [1], fraud costs businesses upwards of $485.6 billion annually, creating an urgent need for stable and effective fraud detection systems. Rule-based frameworks are outdated and cannot keep pace with emerging fraudulent techniques or their growing sophistication.

This paper systematically reviews the contemporary state of payment fraud detection studies that employ AI-based techniques, state-of-the-art solutions, implementation

challenges, and recommendations for improved optimisation. Building on the gaps identified by Marazqah Btoush et al. [5] and Hernandez Aros et al. [8], a total of 124 studies published between 2014 and 2025 have been analysed, with in-depth findings organised into a structured taxonomy. We propose a cost-sensitive detection mathematical formulation based on Le Borgne et al. [6], as well as a discussion of experimental results from emerging technologies, such as federated learning, grounded in the findings of Abadi et al. [10]. All socio-academic challenges are examined in the findings section, while open research directions are presented in the final section.

2 Background and Related Work

Credit card fraud, account takeover, synthetic identity, and money laundering are the most common forms of payment fraud that exploit vulnerabilities in digital payment infrastructures. An extensive business rules framework and hand crafted design of related features were additional shortcomings of earlier detection approaches. These limitations were highlighted in the periodic review of the detection of cyber fraud using machine learning in credit cards, as noted by Marazqah Btoush et al. [5] through references to early detection systems.

Since 2014, AI-based systems have evolved from classical machine learning algorithms such as Random forests and Support Vector Machines to advanced deep learning models, including LSTM, CNN, and GNN. According to Sulaiman et al. [3], the performance analysis of LSTM networks for the credit card fraud emphasized an accuracy of 99.2% and a detection rate of 93.3%. Recent work by Cheng et al. [4] indicates that GNN models offer superior capability as they can capture complex relational patterns.

Moreover, the latest trends in AI for fraud detection lean toward explainability, real-time processing, and ensemble methods that effectively handle concept drift and class imbalance, as demonstrated in the works of numerous authors [13] [14] [2]. Regulatory standards for privacy preservation, such as GDPR and PCI DSS, have strengthened industry practices. The aforementioned authors also highlighted emerging approaches such as federated learning, as conducted by Abadi et al. [10], and blockchain integration, as explored by Ketha and Provodnikova [11]. Quantum Machine Learning (QML) has also been reviewed in various literature reports [12].

3 Methodology

This systematic review strictly followed established protocols aligned with PRISMA guidelines and systematic literature review protocols, which can be drawn from their methodology by Hernandez Aros et al. [8]. A corpus of peer-reviewed articles from leading academic venues such as *IEEE Transactions*, *ACM Computing Surveys*, *Springer* journals, *Elsevier* publications, and high-impact *arXiv* preprints published between 2014 and 2025, was systematically retrieved.

The search strategies used domain specific keywords such as “fraud detection”, “artificial intelligence”, “machine learning”, “deep learning”, “explainable AI”, “real time deployment”, “concept drift”, “class imbalance”, and “optimization”, leveraging Boolean operators and proximity searches.

Identification
Records identified through database searching (n = 2,847)
Screening Records after duplicates removed (n = 1,923)
Eligibility Full-text articles assessed for eligibility (n = 287)
Included Studies included in qualitative synthesis (n = 124)

Fig. 1. PRISMA flow diagram representing the systematic literature selection process.

The inclusion criteria considered: peer reviewed publications in the English language discussing a form of AI/ML application for fraud detection, involving empirical evaluation on real or realistic datasets, published between 2014 and 2025, and demonstrating methodological rigor and reproducibility. The exclusion criteria removed non peer reviewed sources, solely theoretical works without empirical validation, and studies with low methodological disclosure.

The systematic literature selection process is demonstrated in Figure 1, which shows the PRISMA flow diagram tracking the steps from the initial database search to the selection, eligibility evaluating and final inclusion. A total of 2,847 records were digitally retrieved across databases and after duplicate removal, 1,923 records were further screened. Of these, 287 full-text articles were assessed and 124 high-quality studies were included in the qualitative synthesis. These included works vary in geographies, institutions, and methodological approaches.

Figure 1 **Description:** This figure illustrates the PRISMA flow diagram showing the systematic literature selection process. The data progression for each stage is shown with defining figures: *Identification* — 2,847 records were found through database searches; *Screening* — 1,923 records were screened after duplicates were removed; *Eligibility* — 287 full-text articles were assessed for eligibility; and *Included* — 124 studies were added to the final qualitative synthesis.124 studies were added to the final qualitative synthesis.

Each step is arranged in a tabular format with progression arrows (not shown in this simplified tabular layout) to illustrate the filtering procedure that ensured the inclusion of only pertinent and high-quality studies. The quality assessment criteria included methodological rigor, experimental design, dataset representativeness, evaluation appropriateness, and reproducibility. A total of 100 high-quality papers captured diverse regions, institutions, and methodological considerations.

4 Taxonomy of AI-based Approaches

4.1 Supervised Learning Models

A large proportion of the supervised fraud detection approaches from Table 1, which rely on prediction modeling, use labeled transaction data from historical labeled samples. As discussed in the literature where tree ensemble methods are a popular component of fraud

research, there's a clear prevalence of classical algorithms. According to the survey by Hernandez Aros et al. [8], Random Forest is the most used ensemble algorithm, compiled in 74 references (63.8% of the total). Meanwhile, El Bazi Abdelghafour et al. [14] note that for stacking ensemble model studies, XGBoost and LightGBM outperform all other algorithms and other standard datasets with AUC higher than 0.98, meaning these data were compiled with review scores greater than 0.98. **Tree-based Methods:** Random Forest is the most used method (74). Such popularity suggests interpretability and outlier robustness, which is mixed-type agnostic. According to a survey by Hernandez Aros et al. [8]. In the stacking ensemble study by El Bazi Abdelghafour et al. [14], XGBoost and LightGBM outperformed other models, achieving AUC values exceeding 0.98.

Distance Based Methods: K-Nearest Neighbors (KNN) boasts 39 citations and resembles novelty detection as it relies on the closest neighboring points to any transaction. Support Vector Machines (SVM) are noted by Sulaiman et al. [3] for their performance in high dimensional and convoluted characteristics; however, detection effectiveness, if successful, challenges online operational scalability.

Linear Models: At the same time, logistical regression could serve as an easily interpretable baseline model with not many associated computational costs. Unfortunately, limited linearity of expressiveness through feature mapping can be restrictive. At the same time, high-dimensional feature spaces complicate such predictions. Sparse regularization efforts like LASSO, Ridge and Elastic Net help.

4.2 Unsupervised and Anomaly Detection

Unlabeled data-based unsupervised approaches for the detection of zero-day fraud. Frauds are generally incident-specific one, and the known protocols are less susceptible to zero-day fraud predications. They perform very well in situations where attackers change their attacks repeatedly to bypass the detection.

Density-based Methods: Isolation Forests have reported good results in anomaly detection, and its computational complexity is 0(n log n), so it is possible to apply real-time detections. Rezapour [9] applied them to multiple datasets. On the other hand, One Class SVM is pretty robust in creating a boundary for anomaly, but needs some parameter search to get the good performance.

Statistical Methods: The method of robust Mahalanobis distance finds outliers in multivariate distributions. By combining PCA with Mahalanobis distance (called in this context Mahalanobis PCA), you can take a low-rank approximation of the data at once and then apply it to the high-dimensional data efficiently, as considered by Rezapour [9].

Clustering-based Approaches: K-means clustering with outlier detection can help you detect transactions that stand outside typical behavioral clusters. On the other hand, DBSCAN has a good performance in identifying outliers as noise points and processing clusters with various shapes.

4.3 Deep Learning Models

Intelligent representations for fraud detection tasks are made possible by deep learning architectures' ability to recognize temporal and spatial patterns.

Recurrent Neural Networks (RNNs): According to Sulaiman et al. [3]. These models are excellent at learning sequential transaction data, attaining high performance metrics like 99.2% accuracy and a 93.3% detection rate on synthetic datasets. Contextual understanding is further enhanced by bidirectional LSTMs.

Convolutional Neural Networks (CNNs): By representing sequential data as two-dimensional images, CNNs are able to identify patterns in transaction behavior and spatial relationships [3].

Autoencoders: By calculating reconstruction errors, autoencoders can detect anomalies by learning the distribution of typical transactions. When it comes to identifying minute irregularities in transaction patterns, unsupervised variations like Variational Autoencoders (VAEs) have demonstrated encouraging results.

Transformer Architectures: In sequence-to-sequence fraud detection tasks, transformer-based models have demonstrated state-of-the-art performance. Nevertheless, despite their superior contextual learning capabilities, their high computational requirements may restrict real-time applicability.

4.4 Graph-based Methods including Graph Neural Networks

Methods that utilize graphs to create models of fraud networks can show all types of relationship patterns and effects between the parts of a transaction network.

Classic Graph Analytics: Legal Analysis's graph algorithms, such as PageRank, centrality measures, and community detection have assisted Law Enforcement in identifying suspicious account clusters and possible money laundering activities. Various metrics, including graph density and clustering coefficients, have been developed to provide interpretable measures of fraud.

Graph Neural Networks (GNNs): The most comprehensive review of financial fraud detection using GNN was published by Cheng et al. [4], and categorized it into 4 primary categories, including Convolutional, Attention, Metapath and Transformer based architecture, etc. Duan et al. [18] introduced the Causal Temporal Graph Neural Network (CaT-GNN) which achieved state of the art performance due to its ability to model temporal causation within transactional graph structures using GNN architecture. The way in which a GCN works (Graph Convolution Neural Network) is to aggregate a node's own information and perform gradient propagation along the edges connecting neighbouring nodes in order to extract structural dependencies from those nodes' neighbourhoods.

Dynamic Graph Models: These models take into account the dynamic nature of an evolving relationship between nodes over time, which is needed for identifying adaptive manners of conducting fraud by changing the overall structure of the network over time.

4.5 Hybrid and Ensemble Models

The focus of ensemble methods is to combine various learning models in order to build stronger predictors rather than relying on a single model. Stacking Ensembles: In Stacking, performance enhancement is achieved through the combination of multiple base learners, which may include Random Forest, XGBoost, or Neural Networks, in

Table 1. Taxonomy of AI Models in Payment Fraud Detection

Approach	**Representative Algorithms**	**Key Advantages**
Supervised Learning	Random Forest, XGBoost, SVM, KNN	High accuracy, interpretability , robust performance
Unsupervised/Anomaly	Isolation Forest, One-class SVM, Autoencoders	Novel pattern detection, no labeled data required
Deep Learning	LSTM, CNN, Transformers, VAE	Complex pattern modeling, temporal dependencies
Graph-based	GNN, CaT-GNN, GraphSAGE	Network structure exploitation, relational patterns
Hybrid/Ensemble	Stacking, Boosting, Voting	Superior robustness, adaptive capabilities

addition to a meta-learner. Almalki et al. [2] showcased this well in their work by noting that the ensemble of multiple layers which was built achieved an impressive AUC-ROC of 91.8%. The addition of Bayesian model averaging is also notable as it assists with the quantification of uncertainty in the predictions. Boosting Methods: The performance of CatBoost, AdaBoost, and LightGBM exceeds expectations as a result of their innovative iterative combination of weak learning models. CatBoost is particularly powerful in transactional datasets because it natively processes categorical data. Voting Classifiers: In both soft and hard voting ensembles, the predictions of several models are combined and thus provide better accuracy and stability over the predictions of individual learners.

Summary Table: Table 2 presents a comprehensive taxonomy of AI models in the fraud detection domain, categorizing them by methodological approach and application type. This structured taxonomy serves as a useful reference for identifying suitable AI techniques for different types of payment fraud detection tasks.

Description for Table 1: We classified AI-based fraud detection into five main categories in this table to give an overall taxonomy of such algorithms. We will give some typical examples in each category.

Supervised learning techniques: These include high accuracy and interpretability models such as Random Forest, XGBoost, SVM, and KNN. Separating normal behavior from suspicious activities requires large labeled training sets which provide rules to the model.

Unsupervised and anomaly detection approaches: Isolation Forest as well as Autoencoders are effective when the data is not labeled but only when something

different happens from usual. Creating feature representation using stacked convolutional GRU layers further enhances model performance.

Automated Machine Learning (AML): The core benefit of using AML to counter transaction fraud is that financial institutions do not need human operators to identify fraud, auto-rectify possible fraud activities, and save millions in damages or overtime.

Hybrid models: Over time, every organization would likely develop hybrid models that take the good parts of different models and push for deploying them.

5 Real-Time Systems and Deployment Issues

A low latency fraud detection system must process every transaction in less than a millisecond, making accurate decisions while constrained by available processing power. This paper investigates the architectural patterns, deployment mechanisms, and performance tuning considerations for effective implementation and operation of these systems.

5.1 Latency and Throughput Requirements

Today, the lowest required processing latency for authorization decisions in modern payment systems is under 100 ms, with requirements falling below 10 ms for some instant payment networks [7]. Edge IoT platforms can improve latency from 75 ms to 8 ms by employing distributed processing [16].

Throughput requirements are determined by the number of transactions, with large payment processors having to process upwards of 65,000 transactions per second on peak days. Integrating 5G networks increased throughput capacity from 2.5 to 12.8 GB/s, as reported in new studies [7], which ensures real-time fraud detection.

5.2 Edge Computing and Distributed Architectures

Edge computing deployment strategies have been shown to reduce latency by processing transactions near their origin points, i.e., POS terminals and mobile devices. Given models' architectures severed into parts, each processing node at the edge and middle nodes in the cloud, the latency between the two has to be optimized rationally with respect to its impact on detection.

The federated learning framework not only ensures that fraud detection becomes collaborative among financial institutions but also that individual data privacy remains intact, without the possibility of it being leaked. The Starlit framework is shown to scale up federated learning with strong formal security guarantees, preserving client privacy and robustness stability with respect to client outages. This helps address some of the major problems in federated learning, such as lack of formal security definitions, high computationally expensive scalability (O (n2)), and vulnerability to client dropouts.

5.3 Model Optimization and Compression

Model compression techniques can reduce the edge deployment cost. The compression techniques used in this work include pruning, quantization, and knowledge distillation. We obtained state-of-the-art quantized neural networks with 4× memory reduction, very close to the baseline. Quantum optimality and its performance should vary based on trust scores for different transactions, to lighten the load of model deployment for outgoing and incoming transactions. By dynamically selecting models for different transactions, the average processing is reduced by 40% while maintaining high detection accuracy.

5.4 Streaming and Incremental Learning

Streaming machine learning algorithms continuously process transaction flows, gradually updating models to deal with concept drift. In their paper on *Heterogeneous Adaptive Ensemble Learning of Concept Drift: A Case Study with Credit Card Fraud Detection*, Museba and Vanhoof [13] brought stability under distributional shift to light with online random forest and incremental SVM variants.

Efficient neural network updates occur with minibatch gradient descent, which has well-bounded memory requirements. Reservoir sampling maintains a representative subset of training samples so that concept drift triggers model retraining at appropriate times.

6 Optimization and Decision Policies

An ideal detection system should balance several objectives that are fundamentally in conflict with each other, including detection accuracy, false positive rates, processing latency, and operational costs. This section introduces formalisms and strategies for optimization.

6.1 Cost-Sensitive Learning Formulation

The new cost sensitive learning formulation explicitly accounts for the asymmetric impact of false positives and false negatives in a fraud detection context. Formulated in *The Fraud Detection Handbook* by Le Borgne et al. [6], the cost-sensitive risk function is defined as:

$$R_{\text{cost}}(h) = \sum_{(x,y)\in\mathcal{D}} C(y, h(x)) \cdot P(x, y) \tag{1}$$

where $C(y, h(x))$ is the cost matrix:

$$C(y, h(x)) = \begin{cases} 0 & \text{if } y = h(x) \\ C_{FP} & \text{if } y = 0,\ h(x) = 1 \\ C_{FN} & \text{if } y = 1,\ h(x) = 0 \end{cases} \tag{2}$$

The ideal decision threshold τ^* minimizes the expected loss:

$$\tau^* = \arg\min_{\tau} \left[C_{FP} \cdot P(h(x) = 1|y = 0) + C_{FN} \cdot P(h(x) = 0|y = 1)\right] \tag{3}$$

6.2 Multi-Objective Optimization

Fraud detection systems in real life are optimized under the presence of multiple conflicting objectives. The multiobjective formulation considers accuracy, latency, and cost constraints as follows:

$$\begin{aligned} \max_{\theta} \quad & f_1(\theta) = \text{Accuracy}(\theta) \\ \max_{\theta} \quad & f_2(\theta) = \text{Precision}(\theta) \\ \min_{\theta} \quad & f_3(\theta) = \text{Latency}(\theta) \\ \text{s.t.} \quad & \text{Recall}(\theta) \geq \gamma_{\min} \end{aligned} \tag{4}$$

Pareto-optimal solutions define trade-off frontiers, allowing system designers to choose operating points based on business requirements and overall risk appetite.

6.3 Dynamic Threshold Optimization

The introduction of adaptive threshold mechanisms allows for the update of decision boundaries subject to fraud prevalence and operational restrictions. The dynamic threshold at time t, denoted τ_t, is given by:

$$\tau_t = \tau_{t-1} + \alpha \cdot \nabla J(\tau_{t-1}) \tag{5}$$

where the composite objective function $J(\tau)$ intertwines detection performance and operational costs.

6.4 Reinforcement Learning for Policy Optimization

Policy or decision-making for fraud detection can be optimized using reinforcement learning frameworks that learn from interaction outcomes over time. Describing fraud detection as a Markov Decision Process (MDP), we define the state–action–reward formulation as:

$$Q^*(s, a) = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t r_{t+1} \,\middle|\, s_0 = s,\ a_0 = a\right] \tag{6}$$

Here, transaction contexts correspond to states s, detection decisions to actions a, and business outcomes to rewards r.

7 Evaluation Metrics and Practices

7.1 Performance Metrics

Assessing fraud detection models involves the utilization of metrics that consider the imbalance of categories but also factor in the real-world implications of mistakes. The use of traditional classification metrics is still relevant, but they require a shift in

understanding. For example, precision and recall: precision measures the proportion of fraud cases that were correctly identified compared to the total number fraud cases predicted. Recall, however, measures the proportion of actual fraud cases that were successfully identified.

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN} \tag{7}$$

F1-Score: The harmonic mean of precision and recall balances both aspects, especially when dealing with imbalanced datasets:

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \tag{8}$$

AUC-ROC and AUC-PR: The Area Under the Receiver Operating Characteristic (ROC) curve and the Precision–Recall (PR) curve provide threshold-independent evaluation. Among the two, AUC-PR is often more informative for highly imbalanced fraud datasets [15].

7.2 Business-Oriented Metrics

Classification metrics are used to judge how well the model detects objects, whereas business metrics allow you to see the value of your model in the real world. Business metrics take into account the three main perspectives: financial, operational, and user impact.

Cost-Weighted Accuracy: This metric reflects the different monetary values of false negatives and false positives. In general, the cost associated with missing a fraud case is usually far greater than that of issuing a false alarm.

$$\text{Cost-Weighted Accuracy} = \frac{C_{TN} \cdot TN + C_{TP} \cdot TP}{C_{TN} \cdot TN + C_{TP} \cdot TP + C_{FP} \cdot FP + C_{FN} \cdot FN} \tag{9}$$

Alert Fatigue Metrics: These quantify operational strain caused by excessive false positives. The *Alert Quality Index* The relationship between detection efficiency and the work involved in investigations is expressed through the *Alert Quality Index*.

Financial Impact Metrics: People measure the total economic benefits of fraud detection systems by calculating prevented losses from being defrauded, various costs associated with investigating fraud, and experience of the customer using a fraud detection system. Businesses frequently rely on the Return on Investment (ROI) of fraud detection systems to translate model performance into business-like language.

7.3 Temporal Evaluation Strategies

Fraud detection systems work on time-sensitive, evolving environments. Therefore, time is of the essence [13]. They work on moving windows of a constant size, sliding through the available data [13]. The genuine labels show up after the model has predicted them [13]. The assessment keeps a temporal order [13]. No data should leak, and we

Table 2. Evaluation Metrics for Fraud Detection Systems

Metric Category	Specific Metrics	Application
Classification	Precision, Recall, F1-Score, AUC-ROC, AUC-PR	Model comparison, threshold selection
Business Impact	Cost-weighted accuracy, Financial ROI, Alert Quality Index	Operational optimization, business value
Temporal	Sliding window accuracy, Concept drift detection	Adaptation monitoring, maintenance scheduling
Operational	Processing latency, Throughput, Resource utilization	System performance, scalability assessment

present results based on real-life effectiveness [13]. The metrics primarily used in fraud detection systems can be found in Table 2. Classification metrics (Precision, Recall, F1-Score, AUC-ROC and AUC-PR) are primarily used for model performance and threshold value selection assessment [8, 15]. Business impact metrics (Cost-Weighted Accuracy, Financial Return on Investment, Alert Quality Index) concern the operational effectiveness of the fraud detection systems [6]. Time impact metrics (Sliding Window-based, Concept Drift Detection) concern systems' evolution through fraud development [13]. Performance metrics concern Processing Latency for system effectiveness and processing readiness [16]. The metrics observed in Fraud detection systems stem from the development field of preference [15]. Classification metrics (Precision, Recall, F1-Score, AUC-ROC and AUC-PR) are mostly used for performance assessment and threshold value [15]. Business impact (Cost-Weighted Accuracy, Financial Return on Investment, Alert Quality Index) account for the operational and financial realities of the detection systems [6]. Time impact (Sliding Window-based, Concept Drift Detection) impact the system's ability to adapt over time with the evolution of fraud [13]. Finally, the operational impact of processing latency is assessed by the system for effectiveness and operational readiness [16].

8 Open Challenges and Future Directions

8.1 Concept Drift and Adaptive Learning

The trends of fraud are also evolving constantly as fraudsters strive to evade detection. This is termed the concept drift, which poses a challenging problem since it is constant. The performance of the models can deteriorate very rapidly under different distributions of data. There is then a need to have learning strategies that would guarantee accuracy in the long run. Museba and Vanhoof [13] proposed an adaptive heterogeneous ensemble learning credit card fraud detection method and highlighted the importance of dynamism

in adaptation strategies. Such dynamically weighted ensemble techniques have shown greater guarantee to address drift, as opposed to meta-learning techniques which can adapt faster to up-and-coming fraud trends with minimal labelled data. Leverage is also conferred to transfer learning for new condition re-training in related fields to accelerate model re-training.

Future directions are: (1) Theory of drift detection guarantee algorithms, (2) Design of realistic drifting test data patterns, and (3) Balancing towards more stable but flexible online learning practices, enhancing versatility and adaptability.

8.2 Explainable AI and Regulatory Compliance

Interpretable models for fraud detection are becoming increasingly necessary, given the new compliance landscape, e.g. the GDPR "Right to Explanation" and the newly established algorithmic accountability regulations. Thus, regulator authorities and banks remain highly concerned about the auditability of complex deep learning systems [17].

While local explanation techniques (e.g. LIME and SHAP) provide post-hoc interpretability for single-instance decisions, global explanation techniques (e.g. feature importance and decision curves) assist significantly in assessing the model's overall decision-making patterns. Thus, transparent models like decision trees or linear classifiers remain consistently interpretable but often at the expense of lower predictive performance. However, Almalki et al. [2] demonstrate the effectiveness of achieving interpretability alongside high accuracy (99%) through the stacking ensemble framework and SHAP values.

8.3 Privacy-Preserving Collaborative Detection

Fraud detection that is cooperative among banks and other financial institutions can significantly increase detection efficiency; however, legal restrictions related to privacy and the competitive nature of data sharing remain major challenges. Tools such as Secure Aggregation enable data to remain confidential while allowing collaborative model training [6, 16].

Abadi et al. [10] proposed the *Starlit* framework, which offers a privacy-preserving federated learning approach capable of scalable deployment while providing privacy guarantees even under client dropout scenarios through certification mechanisms. The use of federated learning eliminates the necessity of centralizing all data by enabling distributed model training, and differential privacy ensures that data protection is mathematically guaranteed. Secure multi-party computation represents another complementary approach that enables institutions to collaborate analytically while maintaining confidentiality [10].

Possible future research directions include: (1) enhancing communication efficiency and scalability of federated fraud detection systems, (2) developing hybrid architectures that combine on-device learning with secure aggregation mechanisms, and (3) establishing privacy-preserving analytics governance models for cross-institutional collaboration [10].

8.4 Quantum Machine Learning Use Cases

Quantum computing shows strong potential to boost classical machine learning operations which makes it an essential tool for fraud detection systems. The unique properties of quantum systems including superposition and entanglement and decoherence resistance enable developers to create new parallel processing methods and optimization approaches. The classification of complex financial data becomes possible through the use of Quantum Support Vector Machines (QSVM) and Variational Quantum Classifiers (VQC) classifiers. Quantum Neural Networks (QNNs) use quantum parallelism to evaluate multiple hypotheses simultaneously and quantum annealing solves optimization problems related to fraud detection with high efficiency. Quantum computing will deliver exponential speedup for three specific tasks which include anomaly detection and graph-based optimizations and particular probabilistic modeling applications.

The detection of complex fraud patterns becomes more efficient through these techniques which show promise for better performance than classical systems according to recent surveys [12].

8.5 Integration with Blockchain and Smart Contracts

Integration with Blockchain and Smart Contracts. The decentralized tamper-proof structure of blockchain technology enables smart contracts to enhance fraud detection and prevention through AI-based risk models. The unalterable nature of blockchain ledgers enables financial transaction tracking and monitoring and ensures complete transparency and accountability. The integration of smart contracts with AI risk models creates automated fraud detection systems which perform pre-defined actions through transaction holds and alerts and secondary verification processes when rules are breached. The platform achieves maximum security for identity verification and cross institutional data sharing because of its built-in trust and auditability features. The system faces ongoing difficulties when trying to scale up its operations and develop effective fraud detection systems and achieve system interoperability between different networks. The development of lightweight consensus algorithms together with privacy-enhancing data exchange protocols should become the main focus for future research to achieve broad market adoption.

8.6 High-Level Conceptual Framework for Next-Generation Fraud Detection

Based on findings from current investigations [12], we introduce a conceptual model for a next generation fraud detection system combining edge computing, federated learning, explainable AI and dynamic optimization. The model grounded in distributed intelligence is more interpretable and can be adjusted on the fly. Main Features: - Edge Inference Nodes: Ultra-low-latency transaction screening with lightweight detection models situated on the edge of the network. Federated Model Aggregator: Privacy and regulation-compliant federated learning protocol to securely combine model updates. – Explainability Engine: Offers model local (LIME/SHAP) and global interpretability for regulatory clarity and model auditability. – Dynamic Optimizer: Adjust model parameters and decision thresholds on the fly through real-time risk scoring and system

feedback. – Cloud Analytics Hub: Centralized facility for training models, monitoring performance, and concept drift detection using sliding window evaluation tools to trigger monitoring.

8.7 Dataset Limitations and Bias

The availability and quality of datasets remain major limitations in fraud detection research. Public datasets have unrealistic dimensions, diversity and scales in comparison to production-grade environments, leading to incomplete and biased models. Historical data further adds to the temporal lag and outdated models, and fraud evolves, models won't capture the latest patterns and losses exponentially escalate. Undetected and unreported cases of fraud remain excluded from unbalanced datasets leading to selection bias, and models will unrepresent unrefined sophisticated fraud resulting in real-world data undergeneralization. Ground-truth labels generated by manual investigations and legacy systems of fraud detection incorporate biases and tunnel visions novel fraud typologies. Therefore, refining the annotation and using semi-supervised learning will strengthen datasets by improving the reliability, rigor, and robustness of the data and models.

8.8 Evaluation Methodology Concerns

And one of the most criticized aspects of research in this area is when research evaluation fails to take into account what it truly is like to deploy such solutions.

Train-Test Leakage: Using cross-validation with temporal considerations can perform a disservice to your performance by giving away the "future" to the model. Boundary conditions of time are respected, yet broken in most cases.

Metric Choice: Accuracy and AUC-ROC are still the most favored metrics to choose; these metrics downplay a solution's effectiveness on a severely imbalanced dataset, yet business-related metrics are hardly prioritized when assessed in an academic setting.

Gap from Deployment: Where benchmark assessments are concerned, they often fail to acknowledge fraud detection restrictions that exist in the real world — *i.e.*, latency requirements, sometimes assuming all features exist in the wild, and that a model may go stale.

9 Conclusion

From consolidated results of more than 124 AI-based fraud detection studies, it is clear that this field has made significant strides [8].

However, relative progress should be perceived as orthogonal to persistent challenges such as latency, regulatory changes, concept drift adjustment, and explainability, which remain major obstacles to achieving compliance [6].

Interdisciplinary studies are needed for interpretable AI that does not conflict with legal and ethical limitations, relational environments such as graph-based fraud modeling [4], and more appropriate assessment criteria that better reflect the transactional reality of the domain. In addition, learning systems adjusted to fraud evolution and

concept drift [13], secure privacy-preserving multi-stakeholder training systems [10], and the integration of cutting-edge solutions such as quantum machine learning [12] or blockchain technologies [11] are necessary.

A forthcoming complete analysis of that study will be published in the International Journal of Very Big AI Advances.

ACKNOWLEDGMENTS

We extend our gratitude to those organizations part of our academic-industry partnership for their steadfast assistance; on the other hand, most of the banks are notoriously mum bush league.

References

1. Nasdaq, "Global Financial Crime Report 2024," Nasdaq Inc., Tech. Rep., 2024.
   DOI: https://www.nasdaq.com/global-financial-crime-report
   URL: https://www.scribd.com/document/701248208/
2. A. Almalki and M. Masud, "Financial Fraud Detection Using Explainable AI and Stacking Ensemble Models," *arXiv preprint arXiv:2505.10050*, 2024.
   DOI: https://doi.org/10.48550/arXiv.2505.10050
   URL: https://arxiv.org/abs/2505.10050
3. S. S. Sulaiman et al., "Credit Card Fraud Detection Using Improved Deep Learning Models," *Computers, Materials & Continua*, vol. 78, no. 1, pp. 1015–1035, 2024.
   DOI: https://doi.org/10.32604/cmc.2023.046051
   URL: https://www.techscience.com/cmc/v78n1/55404
4. D. Cheng et al., "Graph Neural Networks for Financial Fraud Detection: A Review," *arXiv preprint arXiv:2411.05815*, 2024.
   DOI: https://doi.org/10.48550/arXiv.2411.05815
   URL: https://arxiv.org/abs/2411.05815
5. E. A. L. Marazqah Btoush et al., "A Systematic Review of Literature on Credit Card Cyber Fraud Detection Using Machine and Deep Learning," *PeerJ Computer Science*, vol. 9, e1278, 2023.
   DOI: https://doi.org/10.7717/peerj-cs.1278
   URL: https://pmc.ncbi.nlm.nih.gov/articles/PMC10280638/
6. Y. Le Borgne et al., "Cost-Sensitive Learning," in *Fraud Detection Handbook*, 2022, ch. 6.
   URL: https://fraud-detection-handbook.github.io/fraud-detection-handbook/Chapter_6_ImbalancedLearning/CostSensitive.html
7. F. Author et al., "Real-Time Fraud Detection Using AI-Driven Machine Learning Algorithms," *International Research Journal of Modernization in Engineering, Technology and Science*, vol. 7, no. 2, 2025.
   URL: https://www.irjmets.com/uploadedfiles/paper/issue_1_january_2025/66274/final/fin_irjmets1736942998.pdf
8. L. Hernandez Aros et al., "Financial Fraud Detection Through the Application of Machine Learning Techniques: A Literature Review," *Humanities and Social Sciences Communications (Nature)*, vol. 11, no. 1, 2024.
   DOI: https://doi.org/10.1057/s41599-024-03606-0
   URL: https://www.nature.com/articles/s41599-024-03606-0

9. M. Rezapour, "Anomaly Detection Using Unsupervised Methods: Credit Card Fraud Detection," *International Journal of Advanced Computer Science and Applications*, vol. 10, no. 11, 2019.
DOI: https://doi.org/10.14569/IJACSA.2019.0101101
URL: https://thesai.org/Downloads/Volume10No11/Paper_1-Anomaly_Detection_using_Unsupervised_Methods.pdf
10. A. Abadi et al., "Starlit: Privacy-Preserving Federated Learning to Enhance Financial Fraud Detection," *arXiv preprint arXiv:2401.10765*, 2024.
DOI: https://doi.org/10.48550/arXiv.2401.10765
URL: https://arxiv.org/abs/2401.10765
11. S. Ketha and A. Provodnikova, "Combining Blockchain and AI for Fraud Detection: Building Secure, Transparent, and Sustainable Financial Ecosystems," *Global Business and Innovation Summit*, vol. 7, no. 3, 2024.
DOI: https://doi.org/10.59461/gbis.v7i3.599
URL: https://gbis.ch/index.php/gbis/article/view/599
12. I. Author et al., "A Comprehensive Survey on Quantum Machine Learning Algorithms for Fraud Detection in Financial Sectors," *SSRN Electronic Journal*, 2024.
DOI: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4792054
URL: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4792054
13. T. Museba and K. Vanhoof, "An Adaptive Heterogeneous Ensemble Learning Model for Credit Card Fraud Detection with Concept Drift," *Advances in Science, Technology and Engineering Systems Journal*, vol. 9, no. 3, pp. 1–11, 2024.
DOI: https://doi.org/10.25046/aj090301
URL: https://www.astesj.com/v09/i03/p01/
14. E. B. Abdelghafour et al., "Enhancing Credit Card Fraud Detection Using a Stacking Model with Optimized Feature Selection," *International Journal of Advanced Computer Science and Applications*, vol. 15, no. 10, 2024.
DOI: https://doi.org/10.14569/IJACSA.2024.01510110
URL: https://thesai.org/Downloads/Volume15No10/Paper_110-Enhancing_Credit_Card_Fraud_Detection.pdf
15. Keylabs AI, "Understanding the F1 Score and AUC-ROC Curve," Keylabs AI Blog, 2024.
URL: https://keylabs.ai/blog/f1-score-auc-roc/
16. Transmit Security, "Stop Fraud Sooner — at the Edge," Transmit Security Blog, 2024.
URL: https://transmitsecurity.com/blog/edge-fraud-detection
17. L. Author et al., "Explainable AI (XAI) for Fraud Detection: Building Trust and Transparency," *SSRN Electronic Journal*, 2025.
DOI: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4792054
URL: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4792054
18. Y. Duan et al., "CaT-GNN: Enhancing Credit Card Fraud Detection via Causal Temporal Graph Neural Networks," *arXiv preprint arXiv:2402.14708*, 2024.
DOI: https://doi.org/10.48550/arXiv.2402.14708
URL: https://arxiv.org/abs/2402.14708